

Tracing Heterogeneous APIs (CUDA, OpenCL, L0, OpenMP)

Thomas Applencourt Brice Videau

Argonne National Laboratory

25th Oct 2022

Objective

Understand programming models implementation and usages.

Example:

- How programming models are implemented on top of each other?
 - How OpenMP nowait are implemented in LLVM?
- How applications are using programming models?
 - What is the maximum memory allocated by my program on the GPU?

Solution

- Trace as many programming models as possible
 - Trace should capture as much context as possible, and be lightweight as possible
- Develop tools to analyze traces

Programming-Model Centric Debugging / Tracing

```
18:56:59.677295870 - arc03 - vpid: 37040, vtid: 37040
- lttng_ust_ze:zeKernelSetIndirectAccess_entry:
  { hKernel: 0x0000000002cd2b20, flags: [ ZE_KERNEL_INDIRECT_ACCESS_FLAG_DEVICE ] }
18:56:59.677296042 - arc03 - vpid: 37040, vtid: 37040
- lttng_ust_ze:zeKernelSetIndirectAccess_exit:
  { zeResult: ZE_RESULT_SUCCESS }
```

- Flexible
 - Fine granularity, you can enable/disable individual events tracing,
 - Trace can be read programmatically (C, Python, Ruby),
 - We provide tools calibrated to our needs as starting-blocks.
- Low/Reasonable overhead

THAPI Consist in 2 bigs components

Open source at: <https://github.com/argonne-lcf/THAPI>

- The tracing of events
 - Use low level tracing: Linux Tracing Toolkit Next Generation (LTTng):
 - Well maintained and established (used in industry leading data-centers)
 - Binary format, about 0.2us overhead per tracepoint (in our case)
 - Tracepoints are generated from APIs' headers
- The parsing of the trace
 - Use Babeltrace2 library and tools (reference parser implementation of Common Trace Format)
 - Pretty Printer, Tally, Timeline/Flamegraph, ...

Supported APIs

- OpenCL, Level Zero, Cuda Driver
- OMPT

THAPI Examples: `iprof -t ./a.out`

Wrapping the API entry points to be able to reconstruct the context.

```
> ./iprof -t ./a.out
{ thread_type: ompt_thread_initial, thread_data: 0x00007f5b0cf0ac48 }
ompt_callback_target:
{ kind: ompt_target, endpoint: ompt_scope_end, device_num: 0, task_data: 0x0000000000000000,
  target_id: 1, codeptr_ra: 0x00007f5b26fa47e0 }
[...]
ompt_callback_target_data_op_intel:
{ endpoint: ompt_scope_begin, target_id: 1, host_op_id: 7, optype: ompt_target_data_transfer_to_device,
  src_addr: 0x00007f5b20088280, src_device_num: -10, dest_addr: 0xffffc001ffd80000,
  dest_device_num: 0, bytes: 131072, codeptr_ra: 0x00007f5b26fa47e0 }
clEnqueueMemcpyINTEL_entry:
{ command_queue: 0x181a540, blocking: CL_FALSE,
  dst_ptr: 0xffffc001ffd80000, src_ptr: 0x00007f5b20088280, size: 64, num_events_in_wait_list: 0,
  event_wait_list: 0x0, event: 0x7ffc4ac01378, event_wait_list_vals: [] }
clEnqueueMemcpyINTEL_exit:
{ errcode_ret_val: CL_SUCCESS, event_val: 0x1dffb30 }
ompt_callback_target_data_op_intel:
{ endpoint: ompt_scope_end, target_id: 1, host_op_id: 7, optype: ompt_target_data_transfer_to_device,
  src_addr: 0x00007f5b20088280, src_device_num: -10, dest_addr: 0xffffc001ffd80000,
  dest_device_num: 0, bytes: 131072, codeptr_ra: 0x00007f5b26fa47e0 }
```

THAPI Examples: iprof

```
$iprof ./target_teams_distribute_parallel_do.out # Using Level0 backend
```

```
Trace location: /home/tapplencourt/lttng-traces/iprof-20210408-204629
```

```
BACKEND_OMP | 1 Hostnames | 1 Processes | 1 Threads |
```

Name	Time	Time(%)	Calls	Average	Min	Max
ompt_target	3.65ms	100.00%	1	3.65ms	3.65ms	3.65ms
Total	3.65ms	100.00%	1			

```
BACKEND_OMP_TARGET_OPERATIONS | 1 Hostnames | 1 Processes | 1 Threads |
```

Name	Time	Time(%)	Calls	Average	Min	Max
ompt_target_data_alloc	1.97ms	54.19%	4	491.63us	847ns	1.12ms
ompt_target_data_transfer_to_device	1.26ms	34.63%	5	251.37us	112.60us	460.90us
ompt_target_data_transfer_from_device	250.76us	6.91%	1	250.76us	250.76us	250.76us
ompt_target_submit_intel	155.04us	4.27%	1	155.04us	155.04us	155.04us
[...]						
Total	3.63ms	100.00%	11			

```
BACKEND_ZE | 1 Hostnames | 1 Processes | 1 Threads |
```

Name	Time	Time(%)	Calls	Average	Min	Max
zeModuleCreate	846.26ms	96.89%	1	846.26ms	846.26ms	846.26ms
zeCommandListAppendMemoryCopy	10.73ms	1.23%	12	893.82us	12.96us	5.33ms
[...]						
Total	873.46ms	100.00%	117			

```
Device profiling | 1 Hostnames | 1 Processes | 1 Threads | 1 Devices |
```

Name	Time	Time(%)	Calls	Average	Min	Max
zeMemoryCopy(DM)	64.48us	7.14%	1	64.48us	64.48us	64.48us
__omp_offloading_33_7d35e996_MAIN__19	27.84us	3.08%	1	27.84us	27.84us	27.84us
[...]						
Total	902.72us	100.00%	13			

Timeline visualization

Use perfetto/chrome protobuf trace format




Figure 1: timeline

Iprof is Just a Tool on top of THAPI

- Babeltrace2 is a plugin architecture

```
babeltrace2 --plugin-path=$libdir "$@" \  
            --component=filter.zeinterval.interval \  
            --component=filter.ompinterval.interval \  
            --component=sink.xprof.tally
```

- iprof is just one way of analyzing the trace from THAPI
- Bindings for babeltrace2 exist in Python, Ruby, ...
- So users can write their own plugins (e.g. OTF2 convertor, memory footprint tracker, ...)

Conclusion / Future Work

- Trace all the runtime stack!
- In the process of the v1.0 release (big refractoring of the internal)
- Deploying it on Polaris
- MPI api / HIP support
- If you want to colaborate, don't hesitate!