

THAPI/Iprof: An Introduction

Thomas Applencourt, Brice Videau and many others (including future you?)

September 25, 2025

Table of Contents

1. Rant: You should profile your code
2. Ok, so now you're convinced. Back to iprof
3. What kind of Analysis?
4. Conclusion

Module and Usage:

- `module load thapi`
- `mpirun -n 10000 -- iprof -- ./a.out`

Repo:

- <https://github.com/argonne-lcf/THAPI>
- <https://github.com/argonne-lcf/THAPI-spack>
- <https://docs.alcf.anl.gov/aurora/performance-tools/iprof/>

Rant: You should profile your
code

HPC is the art of removing bottlenecks

- You should know your bottleneck. (MPI? GPU?)
- Optimizing GPU kernels is fun! But more or less useless¹
- Just ask for more nodes, and scale more. We do HPC.

¹It's obviously an exaggeration, but need to keep this entertaining.

You should have a "performance model"

- You should know if you are slow or fast
- Meaning you should know how fast you can be
- "I know I transfer N bytes through the NIC; the bandwidth is X, so it should take Y, but instead it takes Z..."

Everybody should do that

- No matter what programming model you use (oneCCL, MPI, Pytorch, SYCL, Kokkos)
- We are all scientists. If we're here, we should care about performance
- And performance is not an absolute concept, it's relative to the hardware characteristic.

Ok, so now you're convinced.
Back to iprof

Why a Tracer?

We need to understand what's going on in order to solve bugs²

- Why my data-transfer doesn't overlap (H2D + D2H) ?!
- Why OpenMP Mapping take 10 min?!
- Why my SYCL queue in-order have so much submission overhead?!

Or more high level question:

- Am I GPU bound? MPI Bound? Data-transfer bound?
- What is my memory footprint?
- How does my scaling affect my ratio of MPI/GPU times?
- I need a timeline so I can spot any "bubbles" on my GPU execution

²Or to better use the features given us by vendors...

What are we Tracing?

- THAPI: Tracing Heterogeneous API
- We trace library calls (For Aurora: Level Zero, MPI, OpenMP) and analyze them
 - We are "HPC" First: THAPI/iprof should be able to scale with a low overhead.
 - If you find any bugs, or limitation just slack / email me! Not perfect, but at least we try to improve³
- Based on lot a of amazing open-source tools (LTTng, Babeltrace2, Clang, Perfetto, ...)
- Colleen showed you some "sanity check verification" using the OpenMP backend

³"It is the time you have wasted for your rose that makes your rose so important"

Example openmp

```
tapplencourt@chiatta02:~/tmp/tmp/ALCF_Hands_on_HPC_Workshop/programmingModels/OpenMP/demo/cpp> iprof ./03_map  
Success!
```

```
THAPI: Trace location: /home/tapplencourt/thapi-traces/thapi_aggreg--2025-09-24T20:44:40+00:00
```

```
BACKEND_OMP | 1 Hostnames | 1 Processes | 1 Threads |
```

Name	Time	Time(%)	Calls	Average	Min	Max
ompt_callback_target_emi:target	1.31ms	51.63%	1	1.31ms	1.31ms	1.31ms
ompt_callback_target_submit_emi	711.47us	28.07%	1	711.47us	711.47us	711.47us
ompt_callback_target_data_op_emi:transfer_to_device	456.51us	18.01%	2	228.25us	18.03us	438.48us
ompt_callback_target_data_op_emi:transfer_from_device	27.97us	1.10%	1	27.97us	27.97us	27.97us
ompt_callback_target_data_op_emi:alloc	27.37us	1.08%	2	13.69us	914ns	26.46us
ompt_callback_target_data_op_emi:delete	2.94us	0.12%	2	1.47us	693ns	2.25us
Total	2.53ms	100.00%	9			

What kind of Analysis?

What kind of Analysis?

Summary

- Tally / Summary. Give you overview. You should start with that!
 - `mpirun -n 10000 -- iprof -- ./a.out`
 - This will tell you if you spend more time in MPI or on the GPU⁴

⁴Or not there, meaning on the CPU, in this case you should be ashamed. If I see anyone bounded by the Python interpreter they will be required to take a 10h FORTRAN course

Example of MPI

BACKEND_MPI 1 Hostnames 4 Processes 4 Threads							
Name	Time	Time(%)	Calls	Average	Min	Max	
MPI_Barrier	2.00s	44.14%	800	2.50ms	1.83us	124.77ms	
MPI_Init	1.83s	40.34%	4	456.68ms	268.21ms	568.19ms	
MPI_Reduce	686.87ms	15.17%	1600	429.30us	207ns	11.28ms	
MPI_Bcast	8.20ms	0.18%	4	2.05ms	9.25us	6.57ms	
MPI_Finalize	7.30ms	0.16%	4	1.83ms	1.27ms	2.36ms	
MPI_Comm_size	4.46us	0.00%	4	1.12us	1.04us	1.23us	
MPI_Comm_rank	3.10us	0.00%	4	773.75ns	617ns	953ns	
Total	4.53s	100.00%	2420				
Device profiling 1 Hostnames 4 Processes 4 Threads 4 Devices 4 Subdevices							
Name	Time	Time(%)	Calls	Average	Min	Max	
main_l76	1.43min	99.03%	400	214.28ms	208.91ms	224.83ms	
main_l97	459.93ms	0.53%	400	1.15ms	716.16us	5.80ms	
zeCommandListAppendMemoryCopy(M2D)	375.63ms	0.43%	36	10.43ms	80ns	71.97ms	
main_l43	534.08us	0.00%	4	133.52us	17.76us	241.44us	
zeCommandListAppendMemoryCopy(S2M)	110.08us	0.00%	32	3.44us	2.32us	5.84us	
zeCommandListAppendMemoryCopy(M2M)	26.72us	0.00%	12	2.23us	1.84us	2.80us	
zeCommandListAppendMemoryCopy(M2S)	320ns	0.00%	4	80.00ns	80ns	80ns	
Total	1.44min	100.00%	888				

Good, I'm GPU bound! Not lot of data-transfer...Talking about data-transfer

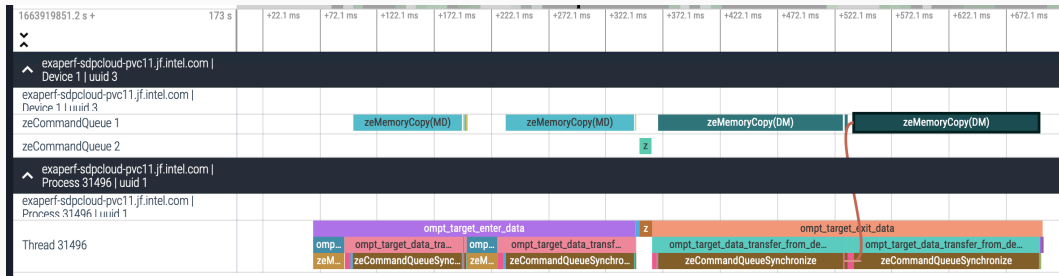
What is slow

- Bandwidth hierarchy: NIC « PCIe « "Xe-Link" GPU « GPU HBM « GPU Cache
- Data-transfer are slow, especially PCIe.
- D == Device Memory, H == Host Memory (pinned), S == Shared Memory (managed), M == "Mallocated" memory

What kind of Analysis?

Timeline

Perfetto Timeline: OpenMP on top of Level Zero




`iprof -l ./a.out` Then open the perfetto in <https://ui.perfetto.dev/>

- Ye showed you some example a few days ago

The timeline backend need a lot of love, if you are interested in visualization just contact me :)

Good and bad code: The bad


Top line: Host Code, Bottom: GPU Code.



Lot of gaps on the GPU -> Bad.

Good and bad code: The "Good"

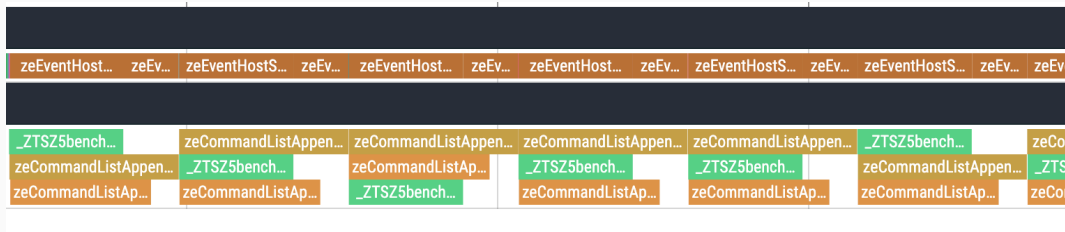
Top line: Host Code, Bottom: GPU Code.



No Gaps... Good?

Good and bad code: The Really Good!

Top line: Host Code, Bottom: GPU Code.



Concurrency! You should always overlap compute and data-transfer.

What kind of Analysis?

Trace

I know some people like to look behind the curtain

Hip on top of Level Zero

```
13:36:02.387547645 - x4204c4s2b0n0 - vpid: 146726, vtid: 146726
- lttng_ust_hip:hipMemset_entry: { dst: 0xff00fffffc4f0000, value: 0, sizeBytes: 12392 }
13:36:02.387550815 - x4204c4s2b0n0 - vpid: 146726, vtid: 146726
- lttng_ust_ze:zeCommandListAppendMemoryFill_entry:
{ hCommandList: 0x0000000004f2da68, ptr: 0xff00fffffc4f0000, pattern: 0x00007fff829294df,
  pattern_size: 1, size: 12392, hSignalEvent: 0x0000000001e672818, numWaitEvents: 2,
  phWaitEvents: 0x0000000001e673d00,
  pattern_vals: "\x00", phWaitEvents_vals: [ 0x0000000001e670658, 0x0000000001ed15bd8 ] }
13:36:02.387558470 - x4204c4s2b0n0 - vpid: 146726, vtid: 146726
- lttng_ust_ze:zeCommandListAppendMemoryFill_exit: { zeResult: ZE_RESULT_SUCCESS }
- lttng_ust_hip:hipMemset_exit: { hipError_t: hipSuccess }
[...]
```

```
iprof -t ./a.out
```

- The more you ask for, the more overhead⁵
- Some analysis are local (tally of tally is a tally), so should be $O(1)$ respectively of the number of hostname⁶
- For now we do only post-processing, so application performance is not impacted. POC for "on-the-fly" is available upon request :)

⁵Duh

⁶Working on doing the same for timeline

What kind of Analysis?

What THAPI/iprof is not

- It's not a full-blown performance analysis framework (Vtune, NSigh, HPC Toolkit, Tau)
- It's not a line-level profiler ⁷

⁷We give you Kernel Time. And we have sampling support of HW counter, but we stop here

Conclusion

Conclusion

- `module load thapi`
- `iprof ./a.out`: so you know where you spend time
- `iprof -t ./a.out`: so you can watch for gap
- `iprof ./a.out`: so you know why

Future work⁸

- Improve timeline visualization
- Improve processing time
- Add new backend (ITT, libfabric)
- Incorporate with other Tracer (python tracer for example)
- Handle multiple binaries

⁸Sound fun, no?