

THAPI/Iprof: Design and Implementation of a Tracer for Heterogeneous API

Thomas Applencourt, Solomon Bekele, Brice Videau and many others

September 25, 2025

Before we start...

- "The audience is very technical", Mike
- So I made the slides in Beamer...

Table of Contents

What?

Why?

Pretty Pictures

How?

Conclusion / Open Questions

One Last note about MPI deployment

What?

What is THAPI / Iprof

- `mpirun -n 60000 -- iprof ./a.out`
- THAPI (Tracing Heterogeneous API) is... a tracer for heterogeneous API.
 - We support OpeCL, LO, Cuda (Driver and Runtime) HIP, OpenMP, MPI
 - We Dump All Arguments. Traces should contain enough information to reconstruct the programming model state.
 - We do hardware counter sampling (frequency, power, traffic...)
- Iprof take the trace, and post-process it (high-level summary, timeline, ...)
- Scalable (tested on 10k GPUs), Low/Reasonable overhead (about 0.2us overhead ns per tracepoint)

What is not?

- It's not a full-blow performance analysis framework (Vtune, NSigh, HPC Toolkit, Tau)
- It's not a line-level profiler ¹

¹We give you Kernel Time. And we know have sampling support of HW counter, but we stop here

Why?

Why a Tracer?

We work on runtime. So we Need to understand what is going on to solve bug.

- Why no data-transfer H2D, D2H overlap?!
- Why OpenMP Mapping take 10 min?!
- Why my SYCL queue in-order have so much submission overhead?!

Application:

- Does I'm GPU bound? MPI Bound? Data-transfer bound?
- What is my memory footprint?
- How does my scaling affect my offload
- And give me some timeline to see if I see some "bubble"

Why a new tracer?

- When starting working on GPU-Aurora, there was not a tracer for Level Zero².
- We wanted just a tracer. Nothing else.
- Want to use "industry standard" binary trace format so we can develop tools on top of it.
- Easy to maintain

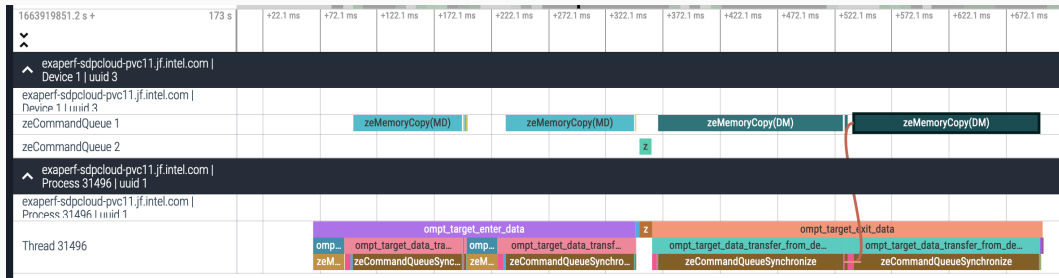
²L0 is the cuda-driver of Intel

Pretty Pictures

Profiling a Simple CUDA Applications with iprof

```
1 > cd /eagle/projects/fallwkshp23/THAPI/CUDA
2 > ./cuda_hello
3 Max error: 0
4 > iprof ./cuda_hello
5 Trace location: /home/videau/litng-traces/iprof-20231010-025849
6
7 BACKEND_CUDA | 1 Hostnames | 1 Processes | 1 Threads |
8
9      Name |      Time | Time(%) | Calls |      Average |      Min |      Max | Error |
10      cuInit | 154.38ms | 43.41% | 1 | 154.38ms | 154.38ms | 154.38ms | 0 |
11      cuCtxSynchronize | 122.67ms | 34.49% | 1 | 122.67ms | 122.67ms | 122.67ms | 0 |
12      cuDevicePrimaryCtxRetain | 71.31ms | 20.05% | 1 | 71.31ms | 71.31ms | 71.31ms | 0 |
13      cuDeviceTotalMem_v2 | 3.01ms | 0.85% | 4 | 753.07us | 731.66us | 772.24us | 0 |
14      cuDeviceGetAttribute | 1.81ms | 0.51% | 392 | 4.61us | 838ns | 172.09us | 0 |
15      cuGetProcAddress | 1.64ms | 0.46% | 373 | 4.39us | 908ns | 531.57us | 0 |
16 [...]
17      cuCtxGetDevice | 2.44us | 0.00% | 1 | 2.44us | 2.44us | 2.44us | 0 |
18      cuDeviceGetCount | 1.47us | 0.00% | 1 | 1.47us | 1.47us | 1.47us | 0 |
19      cuDevicePrimaryCtxRelease | | | 1 | | | | | 1 |
20      cuModuleUnload | | | 1 | | | | | 1 |
21      Total | 355.67ms | 100.00% | 803 | | | | 2 |
22
23 Device profiling | 1 Hostnames | 1 Processes | 1 Threads | 1 Devices | 1 Subdevices |
24
25      Name |      Time | Time(%) | Calls |      Average |      Min |      Max |
26      add | 122.68ms | 100.00% | 1 | 122.68ms | 122.68ms | 122.68ms |
27      Total | 122.68ms | 100.00% | 1 | | | |
```

Perfetto Timeline: OpenMP on top of Level Zero



Simple CUDA Trace

```
1  [...]
2  02:58:48.234853435 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:cuModuleGetFunction_entry: { hfunc: 0x00007ffc1ebf8048,
   ↳ hmod: 0x00000000012e43d0, name: 0x00000000046dcdbd, name_val: "_Z3addiPFS_" }
3  02:58:48.234858185 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:cuModuleGetFunction_exit: { cuResult: CUDA_SUCCESS,
   ↳ hfunc_val: 0x00000000012e1b60 }
4  02:58:48.234867264 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:cuMemAllocManaged_entry: { dptr: 0x00007ffc1ebf8380,
   ↳ bytesize: 4194304, flags: 1 }
5  02:58:48.234900648 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:cuMemAllocManaged_exit: { cuResult: CUDA_SUCCESS,
   ↳ dptr_val: 0x00001471e4000000 }
6  02:58:48.234902604 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:cuMemAllocManaged_entry: { dptr: 0x00007ffc1ebf8378,
   ↳ bytesize: 4194304, flags: 1 }
7  02:58:48.234911613 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:cuMemAllocManaged_exit: { cuResult: CUDA_SUCCESS,
   ↳ dptr_val: 0x00001471e4400000 }
8  02:58:48.239007688 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:cuLaunchKernel_entry: { f: 0x00000000012e1b60, gridDimX:
   ↳ 1, gridDimY: 1, gridDimZ: 1, blockDimX: 1, blockDimY: 1, blockDimZ: 1, sharedMemBytes: 0, hStream: 0x0000000000000000,
   ↳ kernelParams: 0x00007ffc1ebf8300, extra: 0x0000000000000000, extra_vals: [ ] }
9  02:58:48.239034018 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:profiling:event_profiling: { hStart: 0x00000000012e7810,
   ↳ hStop: 0x00000000012e1780 }
10 02:58:48.239035415 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:cuLaunchKernel_exit: { cuResult: CUDA_SUCCESS }
11 02:58:48.239038627 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:cuCtxSynchronize_entry: { }
12 02:58:48.361712445 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:cuCtxSynchronize_exit: { cuResult: CUDA_SUCCESS }
13 02:58:48.370911152 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:cuMemFree_v2_entry: { dptr: 0x00001471e4000000 }
14 02:58:48.371248487 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:cuMemFree_v2_exit: { cuResult: CUDA_SUCCESS }
15 02:58:48.371249954 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:cuMemFree_v2_entry: { dptr: 0x00001471e4400000 }
16 02:58:48.371527993 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:cuMemFree_v2_exit: { cuResult: CUDA_SUCCESS }
17 02:58:48.371538120 - x3015c0s25b0n0 - vpid: 870, vtid: 870 - lttng_ust_cuda:profiling:event_profiling_results: { hStart:
   ↳ 0x00000000012e7810, hStop: 0x00000000012e1780, startStatus: CUDA_SUCCESS, stopStatus: CUDA_SUCCESS, status: CUDA_SUCCESS,
   ↳ milliseconds: 122.68134307861328 }
18 [...]
```

How?

Finally, the Interesting Part!

- With lot of love
- ... and metaprograming.

Two Big Components:

- The tracing of events
 - Using: Linux Tracing Toolkit Next Generation (LTTng) to generate CTF Trace.
 - Tracepoints are generated from APIs' headers
- The parsing of the trace
 - Using: Babeltrace2 ³
 - Babaltrace2 plugin infrastructure generated by MetaBabel
 - Pretty Printer, Tally, Timeline/Flamegraph, ...

iprof is an orchestrator around THAPI, LTTng, and Babeltrace2.

LTTng, and Babeltrace2 are developed by EfficiOS (<https://www.efficios.com/>)

³Similar to ffmpeg pipeline approach

State of the art tracing infrastructure for kernel and user-space.

- Well maintained and established (used in industry leading data-centers)
- Binary format (CTF: Common Trace Format) open standard
- About 0.2us overhead per Tracepoint (in our case: blocking mode)
 - can be relaxed if use case tolerate event losses
- LTTng relay daemons can be setup to stream over the network in complex topologies, or block for "on-the-fly" analysis
 - ideal to deploy at scale
- Fine granularity, you can enable/disable individual tracepoint

Typical Tracer Strategy

To Intersect Symbol:

- we either use `LD_PRELOAD="fake_lib.so" + y dlopen("real_lib.so", RTLD_LAZY | RTLD_LOCAL)`
- or we use Native Support for Interception (OpenCL Layers, OpenMP ⁴)

Implementation of Tracing functions:

```
1 CUresult cuDeviceGetCount(int *count) {
2     tracepoint(lttng_ust_cuda, cuDeviceGetCount_entry, count);
3     CUresult _retval = CU_DEVICE_GET_COUNT_PTR(count);
4     tracepoint(lttng_ust_cuda, cuDeviceGetCount_exit, count, _retval);
5     return _retval;
6 }
```

```
1 21:03:53.070592532 - x3006c0s25b0n0 - vpid: 36056, vtid: 36056
2 - lttng_ust_cuda:cuDeviceGetCount_entry: { count: 0x00007ffe93bec390 }
3 21:03:53.070593929 - x3006c0s25b0n0 - vpid: 36056, vtid: 36056
4 - lttng_ust_cuda:cuDeviceGetCount_exit: { cuResult: CUDA_SUCCESS, count_val: 6 }
```

⁴But not PMPI, we want to trace people who use PMPI...

The Tracepoint Generation seem Tedious

- We trace all APIs entry points (multiple thousand tracepoint...).
- Tedious, error prone, and hard to maintain by hand
- Automatic generation from headers or API description (OpenCL)
 - C99 parser => YAML intermediary representation
 - YAML + user provided meta information + user provided tracepoints => wrapper functions
+ Trace Model
 - Trace Model => tracepoints

Example of Code-gen for LTTNG TracePoint

Cuda Header:

```
CUresult cuDeviceGetCount(int* count);
```

Metadata:

```
1 cuDeviceGetCount:  
2 - [OutScalar, count]
```

```
1 - name: cuDeviceGetCount  
2   type:  
3     kind: custom_type  
4     name: CUresult  
5   params:  
6     - name: count  
7       type:  
8         kind: pointer  
9         type:  
10          kind: int  
11          name: int
```

The Final tracepoints:

```
1 - :name: cuDeviceGetCount_entry  
2   :payload:  
3     - :name: count  
4       :cast_type: int *  
5       :class: unsigned  
6       :class_properties:  
7         :field_value_range: 64  
8         :preferred_display_base: 16
```

```
1 - :name: cuDeviceGetCount_exit  
2   :payload:  
3     - :name: cuResult  
4       :cast_type: CUresult  
5       :class: signed  
6       :class_properties:  
7         :field_value_range: 32  
8       :be_class: CUDA::CUResult  
9     - :name: count_val  
10      :cast_type: int  
11      :class: signed  
12      :class_properties:  
13        :field_value_range: 32
```

And for the reading of the trace with Babeltrace2?

- Thanks you for asking!

Reference parser implementation of Common Trace Format

- Modular plugin infrastructure
- Compose Babeltrace 2 components into trace processing graphs:
 - Sources, Filters, Sinks

```
1 babeltrace2 --plugin-path=$libdir \  
2             --component=filter.zeinterval.interval \  
3             --component=filter.ompinterval.interval \  
4             --component=sink.xprof.tally
```

THAPI use Pipeline of plugins

- Sources are the generated traces
- Filters which aggregate messages
- Sinks which create outputs: Tally, Pretty Print Timeline

Automatic Plugins generation for Babeltrace 2 from the Trace Model

- Problem: Writing Babeltrace 2 plugin by hand is tedious, error prone and hard to maintain.
 - Using Python bindings is too slow -> Use C or C++
- Main Idea: Attaching User-Callbacks to Trace Events
- Metabel generates Babeltrace 2 calls to read, write and dispatch events to User-Callbacks
 - Generate State Machine to handle Babeltrace 2 messages queues
- Open Source: <https://github.com/TApplencourt/metabel>

Code Generated for the Cu Exit

```
1  CUresult cuResult;
2  int count_val;
3  const bt_field *payload_field = bt_event_borrow_payload_field_const(bt_evt);
4  {
5      const bt_field *_field = NULL;
6      _field = bt_field_structure_borrow_member_field_by_index_const(payload_field, 0);
7      cuResult = (CUresult)bt_field_integer_signed_get_value(_field);
8  }
9  {
10     const bt_field *_field = NULL;
11     _field = bt_field_structure_borrow_member_field_by_index_const(payload_field, 1);
12     count_val = (int)bt_field_integer_signed_get_value(_field);
13 }
14 [...]
```

User code:

```
1  #include <metababel/metababel.h>
2  void cuDeviceGetCount_exit_callback(void *btx_handle, CUresult cuResult, int count_val) {
3      std::cout << "cuResult: " << cuResult << ", count_val: " << count_val << std::endl;
4  }
5  void btx_register_usr_callbacks(void *btx_handle) {
6      btx_register_cuDeviceGetCount_exit(btx_handle, &cuDeviceGetCount_exit_callback);
7  }
```


Conclusion / Open Questions

- Scalable Tracer with low-overhead
- Based on industry standard technology (LTTng, Babeltrace2)
- Lot of meta-programming (so maintainable but a small team)
- Should we converge on a meta-data description of API format?

One Last note about MPI
deployment

How to Launch the LTTng Daemon and Post-Processing?

- Problem 1: We need to spawn one daemon per node.
- We need first to spawn the daemon and then launch user-application
- Problem 2: `iprof mpirun ./a.out` Or `mpirun iprof ./a.out`
- I was afraid of doing 1. MPI launcher scare me⁵.

So new Problem... Problem 3:

- Application will call `MPI_init`, but we need a barrier before app run. And We cannot call 2 `MPI_init`
- `MPI_Session_init` "aka" initialize a mpi communicator without call mpi init.

⁵vni, weird argument parsing, ...