

Tracepoint Factory

Automated LTTng Instrumentation & Babeltrace2 Plugin Generation from header files

Thomas Applencourt, Brice Videau | {apl,bvideau}@anl.gov

September 25, 2025

Argonne: Aurora



- “Exascale” Machine
- 10k node, 6 Intel GPUs per nodes
- Developed a Tracing for “Heterogeneous APIs” (OpenCL, Cuda Runtime, Cuda Driver, MPI...) `git@github.com:argonne-lcf/THAPI.git`
- Everything in user space, rely heavily on EfficiOS tools (lttng and babeltrace)


THAPI Example: Summary

| | | | | | | | | | |
|---------------|--|-------------|--|--------------|--|------------|--|---------|--|
| BACKEND_MPI | | 1 Hostnames | | 12 Processes | | 12 Threads | | | |
| Name | | Time | | Time(%) | | Calls | | Average | |
| zeSynchronize | | 2.61h | | 74.75% | | 217672947 | | 43.10us | |
| [...] | | | | | | | | | |
| MPI_Waitall | | 21.86min | | 61.57% | | 87900923 | | 14.92us | |
| MPI_Testall | | 9.94min | | 28.00% | | 87420744 | | 6.82us | |
| [...] | | | | | | | | | |
| Total | | 3.49h | | 100.00% | | 1515216024 | | | |

THAPI Example: Trace (Hip on top of Level Zero)

```
13:36:02.387547645 - x4204c4s2b0n0 - vpid: 146726, vtid: 146726
- lttng_ust_hip:hipMemset_entry: { dst: 0xff00fffffc4f0000, value: 0, sizeBytes: 12392 }
13:36:02.387550815 - x4204c4s2b0n0 - vpid: 146726, vtid: 146726
- lttng_ust_ze:zeCommandListAppendMemoryFill_entry:
{ hCommandList: 0x0000000004f2da68, ptr: 0xff00fffffc4f0000, pattern: 0x00007fff829294df,
  pattern_size: 1, size: 12392, hSignalEvent: 0x000000001e672818, numWaitEvents: 2,
  phWaitEvents: 0x000000001e673d00,
  pattern_vals: "\x00", phWaitEvents_vals: [ 0x000000001e670658, 0x000000001ed15bd8 ] }
13:36:02.387558470 - x4204c4s2b0n0 - vpid: 146726, vtid: 146726
- lttng_ust_ze:zeCommandListAppendMemoryFill_exit: { zeResult: ZE_RESULT_SUCCESS }
- lttng_ust_hip:hipMemset_exit: { hipError_t: hipSuccess }
[...]
```

THAPI Example: Perfetto output



- We usually generate protobuf greater than 2GB... so excited about perfetto presentation tomorrow :) !
- And thanks for Perfetto people fixing our bugs are few years back when we opened on github

- Small Team
- Currently 6 providers, total of 14737 Tracepoints and growing
- 10ich Trace Analysis Plugin

Only one solution: Automation¹!

Take the official header of API and generate:

- The LTTng for tracepoint
- The Babeltrace plugins for analysis infrastructure

We rely heavily on EfficiOS software stack!

¹Or I guess we should vibe-code everything

H2Yaml

How to generation so many tracepoint? From C Header to YAML

- `git@github.com:TApplencourt/h2yaml.git`
- Will generate a YAML file, then from there it's trivial to generate LTTng trace-points²

²Exercise left to the reader, or just talk to Olivier

Example of nice C found in header

```
1 // Forward Declaration
2 typedef struct signal_s signal_t;
3 // Function Pointer with named arguments
4 typedef void (*SignalHandler)(int signum);
5 // Attribute in function call
6 void foo(const signal_t signum, SignalHandler handler);
7 // Declaration, anonymous enum (don't do that...)
8 struct signal_s {
9     enum { S0 } signum;
10 };
```

We will generate

```
functions:  
- name: foo  
  params:  
  - name: signum  
  type:  
    const: true  
    kind: custom_type  
    name: signal_t  
- name: handler  
  type:  
    kind: custom_type  
    name: SignalHandler  
type:  
  kind: void
```

```
enums:  
- members:  
  - name: S0  
    val: 0  
structs:  
- members:  
  - name: signum  
    type:  
      kind: enum  
  name: signal_s
```

```
typedefs:  
- name: signal_t  
  type:  
    kind: struct  
    name: signal_s  
- name: SignalHandler  
  type:  
    kind: pointer  
  type:  
    kind: function  
    params:  
    - name: signum  
      type:  
        kind: int  
  type:  
    kind: void
```

- Using clang, python binding (first proof of concept by EfficiOS Oliver, thanks!)
- Need to deal with lots of idiosyncrasy (anonymous built-in doesn't really work, getting function pointer argument name is atrociously tedious, forward-declared structs point to final node, typedef of struct are parsed twice,)
- But obviously clang can parse everything! So at least their is that...
- h2yaml 100% unit tested, generated same output as our legacy ruby parser.
Can parse all the header we give them now

Example of libclang Madness

```
1  def is_anonymous2(self):
2      match self.kind:
3          case clang.cindex.CursorKind.PARM_DECL:
4              # `is_anonymous()` returns True for `double a` in `void (*a5)(double a, int);`.
5              # We no longer use `not spelling` trick to due to a libclang quirk:
6              # In `(*a6)(a6_t)`, the spelling of `a6_t` will be `a6_t` instead of None.
7              return not self.get_usr()
8          case clang.cindex.CursorKind.FIELD_DECL:
9              # - Unnamed structs have "anonymous ..." in `spelling`
10             # - Named structs within unions: `is_anonymous()` returns True.
11             # - Unnamed bitfields: `is_anonymous()` returns False, but `spelling` is empty.
12             return not self.spelling or "(anonymous at" in self.spelling
13         case clang.cindex.CursorKind.ENUM_DECL:
14             # In `struct S2 { enum { H0 } a; }` where `is_anonymous()` returns False
15             # Fortunately, Clang uses `@EA@` and `@Ea@` in the USR for anonymous enums.
16             # (Though I never saw `@Ea@`...)
17             return self.is_anonymous() or is_in_usr(["@EA@", "@Ea@"])
18     [...]
19     case _:
20         return self.is_anonymous()
```

- Will be nice if multiple tool can agree on the same format
- Will ease the maintenance burden, and avoid duplication of effort
- I know I'm dreaming... But maybe can serve as starting point for your own project.

Metababel

The lesser discussed topic: Analysis of tracez

- We need analyse/transform our ltng trace
- Generate Summary, Timeline...
- We use babeltrace2: we need babeltrace2 plugin
- Python are too slow³ so using C plugin
- But writing Babeltrace C plugin is a little tedious

³Babeltrace2 is too slow too but this is another topic

- Babeltrace2 API is powerfull aka not usable for “common” user.
- Need to understand the CTF hierarchy (common context, environment, stream class) and how to unpack message, push message, handle multiple port, create new class,
- One of our typical plugin is 152066 lines of (generated) babeltrace2 C API

Introducing: Metababel!

- Take a YAML (more or less the one generated by h2yaml), generate all the plugin infrastructure.
- Then one can register callback to particular “event/message” and push new message downstream, without knowing anything about babeltrace2 internal.

Example Metababel YAML format

```
:stream_classes:
- :name: ze
  :event_common_context_field_class:
    :type: structure
    :members:
    - :name: vpid
      :field_class:
        :type: integer_signed
        :field_value_range: 32
        :cast_type: int
    - :name: vtid
      :field_class:
        :type: integer_signed
        :field_value_range: 32
        :cast_type: int
```

```
:event_classes:
- :name: GetPlatformIDs
  :payload_field_class:
    :type: structure
    :members:
    - :name: num_entries
      :field_class:
        :type: integer_unsigned
    - :name: platforms
      :field_class:
        :type: integer_unsigned
    - :name: num_platform
      :field_class:
        :type: integer_unsigned
```

Example: Callback

The only code user need to wrote

```
1
2  #include <metababel/metababel.h>
3  #include <stdio.h>
4
5  static void btx_getplatformids_callbacks(
6      void *btx_handle,
7      void *usr_data,
8      const char* name,
9      int vpid, int vtid, uint64_t num_entries, uint64_t platforms, uint64_t
        ↪ num_platform) {
10
11      printf("Received btx_getplatformids_callbacks message\n");
12  }
13
14  void btx_register_usr_callbacks(void *btx_handle) {
15      btx_register_callbacks_GetPlatformIDs(btx_handle,&btx_getplatformids_callbacks);
16  }
```

Example: Compilation

```
ruby -I../lib ../bin/metabel --component-type SINK \  
    --upstream fake_api.yaml \  
    -o btx_sink  
  
# Wrote your callbacks  
...  
# then link everything together  
gcc -g -o btx_sink.so btx_sink/*.c btx_sink/metabel/*.c  
    -I btx_sink/ -I ../include/ ${CFLAGS}  
babeltrace2 --plugin-path=.  
    --component=source.metabel_source.btx  
    --component=sink.metabel_sink.btx
```

- This is how we handle all our plugins.
- We have the luxury to know statically all our messages types.
- If it's also your case, please give it a shot⁴

Future work

- Improve performance⁵

⁴I'm pretty proud of this project. It served us well!

⁵No more error checking?

Conclusion

- Small Team, need a lot of automation
- h2YAML to generate a yaml from header (so can be used to generate tracepoint and babeltrace plugin)
- Metababel, abstract away babeltrace for ease and speed of trace analysis